

Curiosity-Driven Development of Tool Use Precursors: a Robotic Model

Sébastien Forestier (sebastien.forestier@inria.fr)

Inria Bordeaux Sud-Ouest and Ensta Paristech,
200 Avenue de la Vieille Tour, 33405 Talence, France

Pierre-Yves Oudeyer (pierre-yves.oudeyer@inria.fr)

Inria Bordeaux Sud-Ouest and Ensta Paristech,
200 Avenue de la Vieille Tour, 33405 Talence, France

Abstract

This is the abstract.

Keywords: curiosity-driven learning; tool use; goal babbling; overlapping waves;

Introduction

The ontogenetic development of tool use shows different properties (Guerin, Kruger, & Kraft, 2013): representations are grounded on a large amount of experiences and follow an ongoing process of upgrading [cite others?]. In this paper we focus on one important property in the development of precursors of tool use which is the seamless progression between overlapping phases of behavior with tools and objects, which have been called overlapping waves (Siegler, 1996). The behavioral study of these phases differentiate three categories of behaviors (Guerin et al., 2013). In the first phase, from birth to one year old, babies engage in behaviors without objects, in the second phase, from 4-5 months old, their behavior shifts towards more interaction with a single object, and from 9 months old, begins a simple relational play with several objects. [more details?]

We hypothesize that several mechanisms play a role in the structure of this behavioral progression and in particular 1) the intrinsic motivation to explore as a self-regulation of the learning growth of complexity, and 2) the structure of the representation used to encode sensorimotor experience. Indeed, curiosity-driven learning models with an intrinsic motivation towards situations yielding a high learning progress have shown that developmental trajectories could emerge from the active learning of sensorimotor mappings, in very different settings. In the Playground Experiment (Oudeyer, 2007), a quadruped robot learned how to use its motor primitives to interact with the items of an infant play mat and a robot peer. Also, in a study of the self-organization of vocalizations (Moulin-Frier, Nguyen, & Oudeyer, 2014), an agent had to learn how to use a vocal synthesizer by self-exploration or with the help of humans' demonstrations of phonetic items. In both studies, developmental trajectories of increasing complexity are emerging from learning, with both regularities in developmental steps and diversity. The diversity comes from different mechanisms: random generation in the algorithms, variability in the environment, and the multiples attractors of the dynamic learning system. In those models, the agent learns only one mapping that relates a motor space to a sensory or task space. However, in the perspective of an open-

ended development of reusable skills, and specifically in the development of tool use, multiple interdependent task spaces should be available to the agent as for instance complex actions with tools could make use of previously learned parameterized interaction with the tool.

In this paper we study aspects of those hypothesis leveraging and extending previous models of curiosity-driven learning. We define hierarchies of sensorimotor models that structure the sensory space to reflect the interaction of the different items of the environment. In such hierarchies of models to explore, different exploration choices are available to the agent at each learning iteration: which model to explore, and how to explore that model. This problem is an instance of strategic learning (Nguyen & Oudeyer, 2012), where different outcomes and strategies are available and the agent has to learn which strategies are useful for which outcomes. We define and compare several strategies based on active learning to solve this problem. We compare the different learning conditions in a 2D environment where a simulated arm with three joints plus a gripper can grab tools to move an out-of-reach object. We measure the different phases of behavior during exploration and compare the structure of the behavioral waves. [more results here?]

To our knowledge, this model is the first to show overlapping waves structure in the development of tool use [can we say it is the first model of curiosity-driven development of tool use ?].

However, here we do not study some other important factors in the development of tool use. We consider the hierarchy of models given to the agent and do not study its autonomous building and evolution during developmental stages. Also, social guidance with imitation and mimicry is of central importance for the development of tool use in infants but we do not address the question of its modeling in this paper nor the interplay and tradeoff between social learning and self-exploration. Another important feature of young infants' learning of tool use is the need to adapt to a developing body and to the maturation of vision during ontogenetic development, but here we consider motor control and sensory perception steady over the simulated learning time span.

Other related work. (Ugur, Nagai, Sahin, & Oztop, 2015) (Schmerling, Schillaci, & Hafner, 2015) (Forestier & Oudeyer, 2015) (Sánchez-Fibla, Forestier, Ysard, Moulin-Frier, & Verschure, 2016)

Along with this paper we provide open-source Python

code¹ with iPython/Jupyter notebooks that explain how to reproduce the experiments and analysis.

Methods

Environment

We simulate a 2D robotic arm using tools to move an object into different boxes in the environment. In each trial, we execute a motor trajectory given by the agent, we evaluate its consequences on the sensory dimensions and we give him this sensory feedback. Finally the arm, tools and objects are reset to their initial state. The next sections precisely describe the different items of the environment and their interactions. See Fig.1 for an exemple state of the environment.

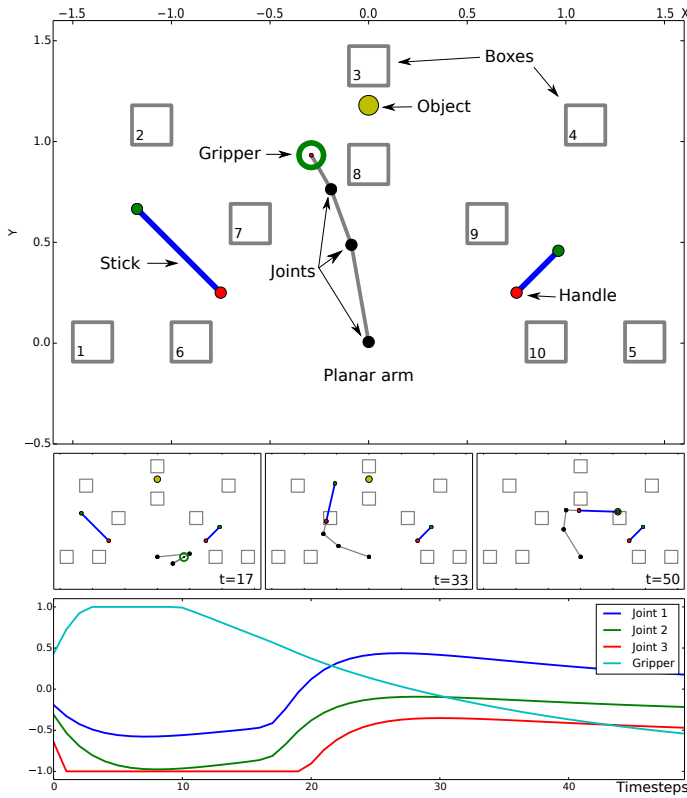


Figure 1: Top: a state of the environment. Middle: position of the arm at time steps 17, 33 and 50 along the 50 steps movement. Bottom: value of the four DMPs during the movement.

Robotic arm The 2D robotic arm has 3 joints plus a gripper located at the end-effector. Each joint can rotate from $-\pi$ rad to π rad around its resting position, mapped to a standard interval of $[-1, 1]$. The length of the 3 segments of the arm are 0.5, 0.3 and 0.2 so the total length of the arm is 1 unit. The resting position of the arm is vertical with each joint at 0 rad and its base is fixed at position $[0, 0]$. The gripper g has 2 possible positions: *open* ($g \geq 0$) and *closed* ($g < 0$) and its resting position is *open* (with $g = 0$). The robotic arm thus

has 4 degrees of freedom represented by a vector in $[-1, 1]^4$. A trajectory of the arm will be represented as a sequence of such vectors.

Motor control We use Dynamical Movement Primitives (Ijspeert, Nakanishi, Hoffmann, Pastor, & Schaal, 2013) to control the arm’s movement as this framework permits the production of a diversity of arm’s trajectories with few parameters. Each of the 4 arm’s degrees-of-freedom (DOF) is controlled by a DMP starting at the rest position of the joint. Each DMP is parameterized by one weight on each of 2 basis functions and one weight specifying the end position of the movement. The weights are bounded in the interval $[-1, 1]$ and allow each joint to fairly cover the interval $[-1, 1]$ during the movement. Each DMP outputs a series of 50 positions that represents a sampling of the trajectory of one joint during the movement. The arm’s movement is thus parameterized with 12 weights, represented by the motor space $M = [-1, 1]^{12}$.

Objects and tools A yellow sphere can be moved into one of ten fixed squared boxes. The initial position of the sphere is $(0, 1.2)$ and is thus unreachable directly with the gripper. Two sticks can be grasped in order to reach the object. A small stick of length 0.3 is located at position $(0.75, 0.25)$ with initial angle $\frac{\pi}{4}$ from the horizontal line. A long stick of length 0.6 is located at position $(-0.75, 0.25)$ with initial angle $\frac{3\pi}{4}$ as in Fig. 1. If the gripper is closed near the handle of one stick (closer than 0.2), it is considered grasped and will follow the gripper’s position and the angle of the arm’s last segment until the gripper opens. Similarly, if the other end of a stick reaches the sphere (within 0.1), the sphere will follow the end of the stick. The ten boxes (identified from 1 to 10) are static and have size 0.2. Boxes 1 to 5 can only be reached with the long stick, and the other five boxes can be reached with both sticks. At the end of the movement, the object is considered to be in one of the box if its center is in the box.

Sensory feedback At the end of the movement, the robot gets sensory feedback from the different items of the environment. It gets the trajectory of the gripper (S_{Hand}), the trajectory of the end of the sticks (S_{Stick_1} and S_{Stick_2}), the position of the object at the end of the movement (S_{Object}), and whether the object is in a box at the end of the movement and the distance between the object and the nearest box (S_{Boxes}). The trajectory of the gripper is represented as the x and y positions and the aperture (1 or -1) of the gripper at 3 time points: steps 17, 33, 50 during the movement of 50 steps (9D). Similarly, the trajectories of the end points of the sticks are 3-points sequences of x and y positions (6D for each stick). The agent receives the identifier (from 1 to 10) of the reached box if one of them has been reached by the sphere, 0 otherwise. He also gets the minimal distance of the object (at the end of the movement) to the center of a box, even if none have been reached. The sensory information thus contains 9

¹Source code and notebooks available as a Github repository at <https://github.com/sebastien-forestier/CogSci2016>

values for the trajectory of the gripper, 6 for the trajectory of the end of each stick, 2 for the end position of the object and 2 for the boxes. The total sensory space S has 25 dimensions.

Learning architectures

The problem settings for the learning agent is to explore its sensorimotor space by iteratively choosing motor parameters that represents arm trajectories and receiving sensory feedback. In this section we describe the different learning architectures that we will compare in the experiment.

Flat architectures We define a flat architecture as directly learning a mapping between the motor space M (12D) and the sensory space S (25D). The control condition is a random motor babbling condition (F-RmB) that randomly chooses new motor parameters m to try at each iteration. In the other conditions, the agent performs Goal Babbling, by self-generating a goal in the sensory space at each iteration and trying to reach it. To do so, the agent needs a sensorimotor model that learns the mapping and provides inverse inference of a probable m to reach a given s . The sensorimotor model stores new information of the form (m, s) with $m \in M$ being the experimented motor parameters and $s \in S$ the associated sensory feedback. It computes the inverse inference with the nearest neighbor algorithm: it gets the motor part of the nearest neighbor of the given s in S , and adds exploration noise (gaussian with $\sigma = 0.01$) to allow new motor parameters to be explored.

The agent also needs an interest model that chooses goals in the sensory space. To generate those goals, different strategies have been studied (Baranes & Oudeyer, 2013). It was shown that estimating the learning progress in different regions of the sensory space and generating the goals where the progress is high leads to a fast learning. However, this idea can't be applied in a 25D sensory space as a learning progress signal can't be properly estimated in this volume. Thus we use a simpler random generation of goals in the sensory space as an interest model in the flat random goal babbling condition (F-RGB). We use the Explauto autonomous exploration library (Moulin-Frier, Rouanet, Oudeyer, & others, 2014) to easily define those sensorimotor and interest models.

Hierarchical architectures The 25D sensory space can be structured to reflect the interaction of the different items of the environment. Indeed, the arm motor position influence the gripper, which influence one of the tools (but not both at the same time), which influence the position of the object and the filling of the boxes. We thus study here learning architectures that could make use of this sensorimotor structure, and we call them hierarchical. Those architecture learn 6 models at the same time (see Fig. ?? : gray squares are models). Each of those models functions in the same way as the random goal babbling flat architecture (F-RGB). Each model has a motor space (e.g. motor space of model 2 is S_{Hand}), a sensory space (respectively S_{Stick_1} , see arrows in Fig. ??, and can choose goals randomly in this sensory space. At each iteration, the

architecture first have to choose the model that will choose a goal, a procedure that we call Model Babbling. Once a model is chosen, it finds a goal in its sensory space, and infer motor parameters (that can be in a sensory space) to reach that goal. Then, it passes those parameters as a goal to be reached by the lower-level model, which similarly infers motor parameters and passes those ones until the actual *Arm* motor space gets parameters to try in the environment (with the same exploration noise as in Flat architectures). Model 4 have also to choose which lower-level model to use in order to reach an object end position s_o in S_{Object} , as two models (3 and 6) have S_{Object} as sensory space. Model 4 chooses the tool that have allowed to reached s_o as close as possible in the past. Finally, when motor parameters m have been tested in the environment and feedback s received, the mappings of all models are updated, but if model 4 chose one tool then the mapping of the other tool is not.

A first condition is to randomly choose the model that will find a goal, this is Random Model Babbling (H-RGB-RMB). The problem of Model Babbling is an instance of strategic learning (Nguyen & Oudeyer, 2012), where different outcomes and strategies to learn them are available and the agent learns which strategies are useful for which outcomes. In that paper, they show that an active choice of the outcomes and strategies based on the learning progress on each of them increase learning efficiency compared to random choice. To develop active learning strategies, we first define a measure of learning progress for each of the 6 models. When a model has been chosen to babble, it draws a random goal s_g , and finds motor parameters m to reach it using the lower-level models. The actual outcome s in the sensory space of the model, associated to m might be very different from s_g as this goal might be unreachable, or because lower-level models are not mature enough for that goal. We define the competence associated to a goal s_g as $C(s_g) = -||s_g - s||$, and the interest $I(s_g)$ associated to this goal as the absolute difference between $C(s_g)$ and the mean competence of the ($k = 20$) nearest previous goals. The interest of a model is initialized at 0 and updated to follow the interest of the goals:

$$I_{model} = \frac{n-1}{n} I_{model} + \frac{1}{n} I(s_g) \quad (1)$$

In condition H-RGB-P-AMB, the choice of model is proportional to their interest (probabilities are proportional to interest but with $\epsilon = 10\%$ of random choice). In condition H-RGB-GR-AMB, the choice of model is greedy (model with maximum interest) but also with $\epsilon = 10\%$ of random choice. Finally, condition H-RGB-P-AMB-PGITC is the same as H-RGB-P-AMB but the choice of the tool to use (model 3 or 6) is with probabilities proportional to the interest of the two models, instead of being based on the more competent tool for the given object goal position.

Results

We perform 100 simulations per condition. Each simulation starts with 100 iterations of motor babbling and then runs for

100000 iterations of the condition. In this section we provide results for different types of measures. First we show an example of a run of a hierarchical condition to illustrate a possible evolution of the interest of each model, and the result in the exploration of the object space. Then, we define a behavioral measure to categorize the different types of behaviors with objects and study the structure of their evolution. We also give a measure of the total exploration of the different spaces during simulations. Finally, we compare the structure of tool choice made to reach object goal position during exploration in the two conditions for which only this choice differs.

Fig. 2 shows details about one trial of the condition H-RGB-P-AMB. We can see the interest of each model during the whole experiment. The interests of models 2 and 5 increase abruptly once the arm succeeded in grabbing the corresponding stick. Following that, the interests of models 3 and 6 also increase abruptly when the object has been touched by the corresponding stick. An example of exploration of the 2D space with the object is also provided in Fig. 2(b) corresponding to the same condition.

We provide a measure of the different types of behaviors with the sticks and the object during exploration. We categorize the behaviors into three types. In the first category (*hand*) are movements of the arm that did not grab any stick and thus not moved the out-of-reached object. The second category (*stick*) are movements that did grab one of the two sticks but did not touch the object with it. The third category (*object*) contains the movements where both a stick was grabbed and the object was moved by the stick. Fig. 3 shows the prototypical evolution of the proportion of the three categories of behavior along the 100000 iterations for conditions H-RGB-GR-AMB and H-RGB-P-AMB. We performed a more detailed analysis by counting the trials where the evolution of the three types of behaviors were similar to the one of Fig. 3(b). A structure was considered similar if it validated each of the following criteria: behaviors of categories *stick* and *object* increase quickly from 0 to more than 10% (potentially after an initial phase with a steady low value), and are followed by a plateau (steady curve with small slope) with no abrupt changes, and behaviors of category *object* start to raise at least 1000 iterations after *stick* started to raise. We counted no structures of that type in Flat conditions, 7 for H-RGB-GR-AMB, 60 for H-RGB-RMB, 70 for H-RGB-P-AMB and 79 for H-RGB-P-AMB-PGITC. Also, the median number of abrupt changes (steady change of more than 10%) in conditions F-RmB, F-RGB, H-RGB-RMB, H-RGB-P-AMB, H-RGB-GR-AMB and H-RGB-P-AMB-PGITC was respectively 0, 1, 2, 2, 6, 1, with a significant difference between condition H-RGB-GR-AMB and the others (Mann-Whitney U tests, $p < 0.0001$).

Also, for each condition we measure the total exploration of the different sensory spaces during training. The exploration of the hand, sticks and object spaces is defined as the number of reached cells in a 100×100 discretization of the




Figure 3: Behavioral measure. Left: H-RGB-GR-AMB, Right: H-RGB-P-AMB.

(X,Y) space of the position at timestep 50 (end of movement). The exploration of the boxes is the number of boxes that have been filled with the object during training. Fig. 4 shows the total exploration of the different sensory spaces for each condition. We provide statistical Mann-Whitney U test results of comparisons of the exploration in different pairs of conditions. Firstly, the Motor Babbling condition (F-RmB) have more explored S_{Hand} and less S_{Object} and S_{Boxes} compared to the other conditions ($p < 0.0001$). Then, F-RGB explores all spaces less than H-RGB-RMB condition ($p < 0.01$). Also, H-RGB-GR-AMB shows lower exploration all spaces than H-RGB-P-AMB ($p < 0.01$). Condition H-RGB-P-AMB-PGITC explores more S_{Stick_2} ($p < 0.05$) than condition H-RGB-P-AMB, and difference is not significant in other spaces.

Fig. 5 shows a comparison of the choice of tool to reach a given object goal position in the conditions H-RGB-P-AMB and H-RGB-P-AMB-PGITC. In those conditions, model 4 learns a mapping between S_{Object} and S_{Boxes} . When this model is babbling, it chooses a random goal $s_b \in S_{Boxes}$ and infers the best object position s_o to reach s_b . To reach s_o , one of the tools (*Stick₁* with model 3 or *Stick₂* with model 6) is chosen. We plot all those choices, at position s_o on a 2D space, with color blue if *Stick₁* was chosen and red if *Stick₂* was chosen, in one figure for each of the two conditions. We can see two very different choice structures. However, goal that can be reached with both tools are more often chosen to be explored with the long stick in the interest-based choice of condition H-RGB-P-AMB-PGITC than in competence-based choice of condition H-RGB-P-AMB.




Figure 5: Chosen tool depending on object goal position. Blue points: long stick choice. Red points: small stick choice. Left: H-RGB-P-AMB, Right: H-RGB-P-AMB-PGITC.

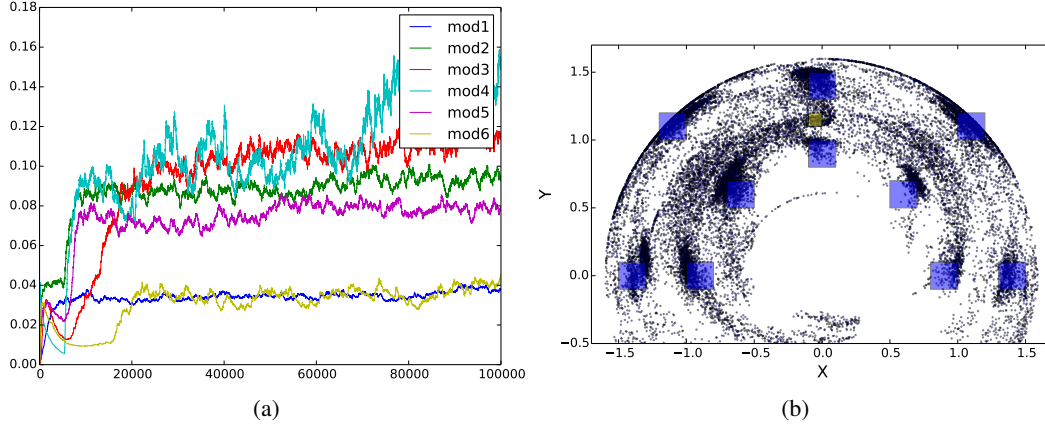


Figure 2: Condition H-RGB-P-AMB. (a) Interests of each model. (b) Exploration of the object space: each dot is one point reached with the object at the end of one movement.




Figure 4: Exploration of sensory spaces.

Discussion

F-RGB vs H-RGB-P-AMB vs H-RGB-GR-AMB

Results show different structures of behavior evolution in those 3 conditions.

H-RGB-P-AMB vs H-RGB-P-AMB-PGITC

Different structure of tool choice: overlapping waves of strategies to fulfill given goals in condition H-RGB-P-AMB-PGITC.

Acknowledgments

References

- Baranes, A., & Oudeyer, P.-Y. (2013, January). Active learning of inverse models with intrinsically motivated goal exploration in robots. *Robotics and Autonomous Systems*, 61(1).
- Forestier, S., & Oudeyer, P.-Y. (2015). Towards hierarchical curiosity-driven exploration of sensorimotor models. In

2015 joint IEEE international conference on development and learning and epigenetic robotics (ICDL-Epirob).

Guerin, F., Kruger, N., & Kraft, D. (2013). A survey of the ontogeny of tool use: from sensorimotor experience to planning. *Autonomous Mental Development, IEEE Transactions on*, 5(1).

Ijspeert, A. J., Nakanishi, J., Hoffmann, H., Pastor, P., & Schaal, S. (2013). Dynamical movement primitives: learning attractor models for motor behaviors. *Neural computation*, 25(2).

Moulin-Frier, C., Nguyen, S. M., & Oudeyer, P.-Y. (2014). Self-organization of early vocal development in infants and machines: the role of intrinsic motivation. *Frontiers in Psychology*, 4.

Moulin-Frier, C., Rouanet, P., Oudeyer, P.-Y., & others. (2014). Explauto: an open-source Python library to study autonomous exploration in developmental robotics. In *ICDL-Epirob-International Conference on Development and Learning, Epirob*.

- Nguyen, S., & Oudeyer, P.-Y. (2012). Active choice of teachers, learning strategies and goals for a socially guided intrinsic motivation learner. *Paladyn*, 3(3).
- Oudeyer, P.-Y. (2007). What is intrinsic motivation? A typology of computational approaches. *Frontiers in Neuro-robotics*, 1.
- Sánchez-Fibla, M., Forestier, S., Ysard, J., Moulin-Frier, C., & Verschure, P. (2016). Unifying affordance and kinematics learning: a computational approach to bimanual affordances. *submitted*.
- Schmerling, M., Schillaci, G., & Hafner, V. V. (2015). Goal-directed learning of hand-eye coordination in a humanoid robot. In *5th joint ieee international conferences on development and learning and on epigenetic robotics (icdl-epirob)*.
- Siegler, R. S. (1996). *Emerging minds: The process of change in children's thinking*. Oxford University Press.
- Ugur, E., Nagai, Y., Sahin, E., & Oztop, E. (2015, June). Staged development of robot skills: Behavior formation, affordance learning and imitation with motionese. *Autonomous Mental Development, IEEE Transactions on*, 7(2).